



CLOUD COMPUTING LAB (22ECAC202)

Activity Report

BATCH: F1

ROLL NO: 131

INSTRUCTOR: Ms. : Sadaf Anjum Savanur

COURSE DETAILS:

Course Name	Cloud Computing
Course Code	24ECAE317
Semester	VI
Division	F

Team Members:

Name	USN	ROLL NO
Bhoomika Marigoudar	01FE22BCI035	131

TOPICS	PAGE NO.
INTRODUCTION AND COURSE OUTCOMES	3
IMPLEMENTATION OF CLOUD SERVICE MODELS (IaaS, PaaS, SaaS)	4
VIRTUALIZATION AND CONTAINERIZATION TECHNIQUES FOR RESOURCE MANAGEMENT IN A CLOUD ENVIRONMENT	6
IMPLEMENT DIFFERENT LEVELS OF AUTOMATION AND CONTAINER ORCHESTRATION USING KUBERNETES	7
MICROSERVICES	9
ANSIBLE	12

Introduction and About Cloud Computing:

Cloud computing is a revolutionary technology that enables on-demand access to a shared pool of configurable computing resources, such as servers, storage, databases, networking, software, and more, over the internet. It allows organizations and developers to scale applications efficiently, reduce infrastructure costs, and accelerate deployment cycles. Through various models like Infrastructure as a Service (IaaS), Platform as a Service (PaaS), and Software as a Service (SaaS), cloud computing provides flexibility and agility in managing IT services. This course introduces core concepts of cloud computing, emphasizing hands-on experience with virtualization, containerization, Kubernetes orchestration, serverless computing, microservices, and DevOps tools, making students industry-ready for modern cloud-based development and operations.

Course Outcomes Overview:

The projects in this course comprehensively address the key aspects of cloud computing. We begin by understanding and applying various cloud services and deployment models to build cloud-based applications. We then explore virtualization and containerization, managing resources efficiently in cloud environments. With Kubernetes, we implement automation and orchestration, streamlining deployment processes. Projects further involve developing scalable applications using serverless architecture and microservices, ensuring efficient event processing and communication. Finally, through tools like Ansible, Puppet, and Chef, we implement DevOps practices to automate configuration management and optimize infrastructure provisioning. Together, these projects align with the course outcomes, providing a holistic learning experience in cloud technologies and practices.

Lab Evaluation 1: Introduction to Cloud Services and Deployment Models

Objective:

- Understand cloud computing concepts, service models (IaaS, PaaS, SaaS).
- Deploy a simple cloud-based application using one of the service models.

Tasks Performed:

- Explored cloud platforms (AWS, Vercel, GCP) and their service offerings.
- Deployed a basic web application on AWS (PaaS).
- Configured environment variables and deployment parameters.

Tools/Technologies Used:

- GCP: Google Cloud ; Cloud Computing Services
- AWS Management Console
- Simple Web App (Node.js,HTML,CSS, React)

Outcome:

- Demonstrated ability to select appropriate cloud service models.
- Successfully deployed a web app on a cloud platform.
- Gained familiarity with cloud deployment processes.

Different Cloud Services Available:

1. IaaS (Infrastructure as a Service)

Definition:

IaaS provides virtualized computing resources over the internet. This includes servers, storage, networking, and operating systems. Users can install their own applications and manage the entire software stack.

Example

AWS EC2, Google Compute Engine, Microsoft Azure VM

Providers:

Use

If a developer wants full control over their virtual machines and network, IaaS is the best choice.

Case:

User

You handle the OS, applications, runtime, data, etc. The cloud provider only manages the infrastructure.

Responsibility:

2. PaaS (Platform as a Service)

Definition:

PaaS provides a ready-to-use platform with an environment for developing, testing, deploying, and maintaining applications. You don't have to manage infrastructure.

Example

AWS Elastic Beanstalk, Google App Engine, **Vercel**, Heroku

Providers:

Use

If you're a developer focusing only on app development and not server setup or OS management, PaaS is ideal.

Case:

User

You write and deploy your code. The provider manages everything else (OS, middleware, servers).

Responsibility:

3. SaaS (Software as a Service)

Definition:

SaaS delivers software applications over the internet on a subscription or free basis. Everything is managed by the service provider.

Example

Gmail, Google Docs, Microsoft 365, Dropbox

Providers:

Use

End-users who want ready-to-use software without worrying about installation or updates.

Case:

User

Only to use the application.

Responsibility:

Vercel Deployment – Step-by-Step Implementation

In our project, we used **Vercel**, which is a **PaaS** platform specialized in deploying **frontend frameworks and static sites**.

Pre-requisites:

- A GitHub account (to store code)
- A Vercel account (linked with GitHub)
- A web app built using **HTML, CSS, Node.js, and React**

Step 1: Push Project to GitHub

1. Create a GitHub repository.
2. Add your local project files (React + Node.js code) to the repo.
3. Commit and push the code:

```
git init
```

```
git add .
```

```
git commit -m "initial commit"
```

```
git remote add origin <repo-link>
```

```
git push -u origin main
```

Step 2: Login to Vercel and Import Project

1. Go to <https://vercel.com> and log in with your GitHub account.
 2. Click on “Add New Project”.
 3. Select the GitHub repository that contains your project.
-

Step 3: Configure Project Settings

1. Vercel will auto-detect the framework (e.g., React).
 2. Set the **build command** (e.g., npm run build) and **output directory** (usually build or dist).
 3. Set **Environment Variables** (if any) like API URLs, tokens, etc.
-

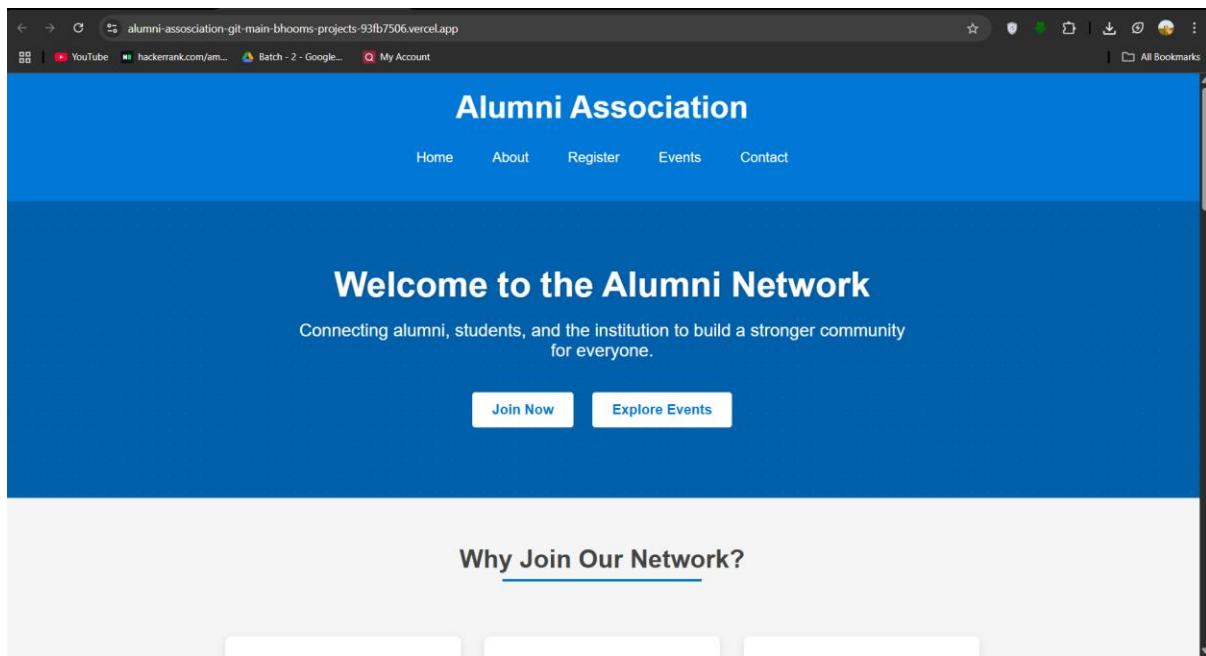
Step 4: Deploy the Application

1. Click on “Deploy”.
 2. Vercel will install dependencies, build the app, and host it.
 3. After a successful deployment, Vercel will provide a live URL (e.g., <https://myapp.vercel.app>).
-

Step 5: Post-Deployment Configuration

- Monitor logs in the Vercel dashboard.
- Set up custom domains (if needed).
- Configure automatic redeploys on GitHub pushes.

Link for Website: <https://alumni-assosciation-git-main-bhooms-projects-93fb7506.vercel.app/>



Lab Evaluation 2: Virtualization and Containerization Techniques

Objective:

- Learn virtualization basics and containerization using Docker.
- Create and manage containers for resource management.

Tasks Performed:

- Installed Docker on local machine.
- Created Docker images and containers using Dockerfiles.
- Managed containers lifecycle (start, stop, restart).
- Explored isolation and resource sharing in containers.

Tools/Technologies Used:

- Docker Engine and CLI
- Docker Hub for images

Outcome:

- Understood virtualization vs containerization differences.
- Created and deployed containers effectively.
- Recognized benefits of containers for cloud environments.

The screenshot shows a web application titled "Latest Mobile Showcase" running on localhost:8080. The page features a header with dropdown menus for "All Brands" and "All Budgets". Below the header are four product cards:

- Samsung Galaxy S23**
Camera: 50MP Triple
Processor: Snapdragon 8 Gen 2
₹74,999
10% Off
- Samsung Galaxy A23**
Camera: 30MP Triple
Processor: Snapdragon 6 Gen 2
₹23,999
5% Off
- iPhone 15**
Camera: 48MP Dual
Processor: A16 Bionic
₹89,999
5% Off
- Redmi Note 13 Pro**
Camera: 200MP Triple
Processor: Snapdragon 7s Gen 2
₹24,999
15% Off

Below these cards is a single card for the **OnePlus 12R**:

- OnePlus 12R**
Camera: 50MP Triple
Processor: Snapdragon 8+ Gen 1
₹39,999
8% Off

At the bottom of the page are navigation buttons: "Prev", "1", "2", "3", and "Next".

The screenshot shows the Docker Desktop interface. The sidebar on the left includes sections for "Ask Gordon" (Beta), "Containers" (selected), "Images", "Volumes", "Builds", "Docker Hub", "Docker Scout", and "Extensions".

The main area is titled "Containers" and displays the following information:

- Container CPU usage: 0.00% / 800% (8 CPUs available)
- Container memory usage: 153.33MB / 2.79GB
- Show charts

A search bar and a filter button ("Only show running containers") are also present.

The container list table has columns: Name, Container ID, Image, Port(s), CPU (%), Last stat, and Actions. The table shows five items:

	Name	Container ID	Image	Port(s)	CPU (%)	Last stat	Actions
1	magical_greider	95ff15ebb9c2	todo-app	8888:80	0%	25 days	D ⋮ T
2	minikube	914a67cc3399	k8s-minikul	0:22	0%	25 days	D ⋮ T
3	mobile-showcas	9a04799ee486	mobile-sho	8080:80	0%	1 minute	D ⋮ T
4	eager_ganguly	c22c6d83a07f	shopping_u	8080:80	0%	1 month	D ⋮ T
5	hospital-manag	-	-	-	0%	6 minutes	D ⋮ T

At the bottom, status bars indicate "Engine running", "Kubernetes running", "Terminal", and "New version available".

Lab Evaluation 3: Kubernetes Orchestration and Automation

Objective:

- Understand container orchestration and automation using Kubernetes.
- Deploy multi-container applications and **manage scaling**.

Tasks Performed:

- Set up a Kubernetes cluster using Minikube.
- Deployed sample multi-container apps with Kubernetes manifests (YAML).
- Explored pods, services, deployments, and scaling.
- Implemented basic automation for rolling updates.

Tools/Technologies Used:

- Minikube (local Kubernetes)
- kubectl CLI
- Sample microservices app (multi-container)

Outcome:

- Learned Kubernetes architecture and components.
- Successfully deployed and managed containerized apps with orchestration.
- Gained skills in automating deployment and scaling processes.

Lab Evaluation 4: Development of Microservices-Based Applications

Hospital Management System – Monitoring with Prometheus and Grafana

Overview

In the Hospital Management System, we employed microservice architecture to break down functionality into separate services:

- patient-service
- doctor-service
- billing-service
- frontend

To monitor the health and performance of these services, we used:

- Prometheus (for collecting metrics)
 - Grafana (for visualizing those metrics)
-

Microservices Used

Each of the microservices was containerized using Docker and run using Docker Compose. They included:

Service Name	Port	Technology
patient-service	5004	Spring Boot
doctor-service	5002	Node.js
billing-service	5003	Spring Boot
frontend	80/3000	React + Nginx

Prometheus Integration

◊ What is Prometheus?

Prometheus is an open-source monitoring tool that collects metrics from configured targets at given intervals, evaluates rule expressions, displays results, and can trigger alerts.

◊ Prometheus Configuration (prometheus.yml)

yaml

scrape_configs:

```
- job_name: 'patient-service'  
  metrics_path: /actuator/prometheus  
  static_configs:  
    - targets: ['patient-service:5004']
```

```
- job_name: 'doctor-service'  
  metrics_path: /metrics  
  static_configs:  
    - targets: ['doctor-service:5002']
```

```
- job_name: 'billing-service'  
  metrics_path: /actuator/prometheus  
  static_configs:  
    - targets: ['billing-service:5003']
```

This configuration:

- **Scrapes metrics from each service every 15 seconds (default).**

- Uses the appropriate metrics endpoint:
 - /actuator/prometheus for Spring Boot (via Micrometer)
 - /metrics for Node.js (via prom-client)
-

With this, the services expose a Prometheus-compatible endpoint at:

bash

CopyEdit

`http://localhost:5004/actuator/prometheus`

Configuring Node.js Services

For doctor-service (Node.js), we used the prom-client library.

Installation:

`npm install prom-client express`

```
const express = require('express');
const client = require('prom-client');
const app = express();
```

```
client.collectDefaultMetrics(); // Enables basic system metrics
```

```
app.get('/metrics', async (req, res) => {
  res.set('Content-Type', client.register.contentType);
  res.end(await client.register.metrics());
});
```

```
app.listen(5002, () => console.log("Doctor Service running on port 5002"));
```

This exposes:

bash

CopyEdit

<http://localhost:5002/metrics>

Common Prometheus Queries Used

Query	Description
<code>up</code>	Displays the status of all targets (1 = UP, 0 = DOWN)
<code>http_server_requests_seconds_count</code>	Shows number of HTTP requests handled by Spring Boot services
<code>http_server_requests_seconds_count{status=~"5..")}</code>	Tracks how many 5xx server errors occurred
<code>rate(http_server_requests_seconds_count[1m])</code>	Rate of HTTP requests per second in the last minute

Grafana Integration

◊ What is Grafana?

Grafana is a powerful visualization tool that allows users to create interactive dashboards from various data sources, including Prometheus.

◊ Grafana Setup in Docker Compose

yaml

grafana:

image: grafana/grafana:latest

container_name: grafana

ports:

- "3005:3000"

environment:

- GF_SECURITY_ADMIN_USER=admin

- GF_SECURITY_ADMIN_PASSWORD=admin

depends_on:

- prometheus

networks:

- monitoring-net

◊ Steps to Visualize Metrics

1. Visit Grafana at <http://localhost:3005>

2. Log in with:

◦ Username: admin

◦ Password: admin

3. Add Prometheus as a data source (URL: <http://prometheus:9090>)

4. Create a dashboard with panels for:

- CPU usage
 - Memory usage
 - HTTP error rates
 - Service uptime
-

🔍 Debugging: Prometheus Not Seeing Services

If up = 0 for any service:

- Check Docker container logs (`docker logs <container-name>`)
 - Ensure metrics endpoint is reachable from Prometheus container
 - Confirm ports and paths match in `prometheus.yml`
 - Visit Prometheus /targets UI: <http://localhost:9090/targets>
-

📦 Final Docker Compose Overview

`yaml`

`services:`

`patient-service: ...`

`doctor-service: ...`

`billing-service: ...`

`frontend: ...`

`prometheus:`

`image: prom/prometheus`

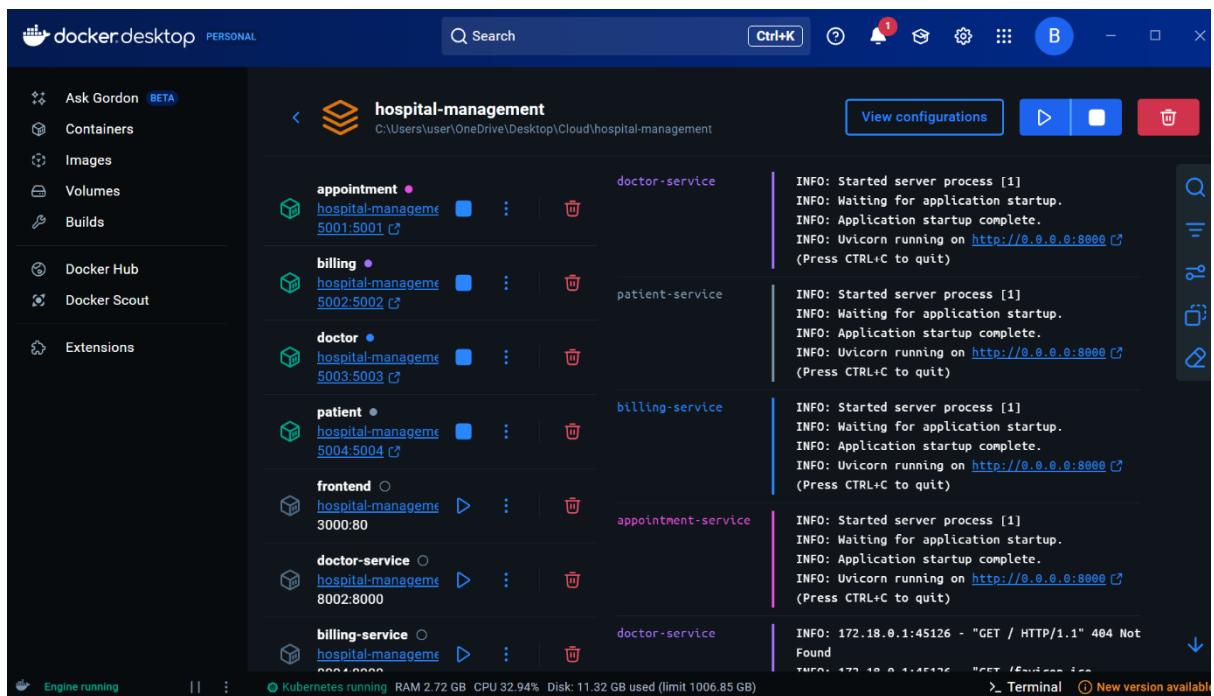
...

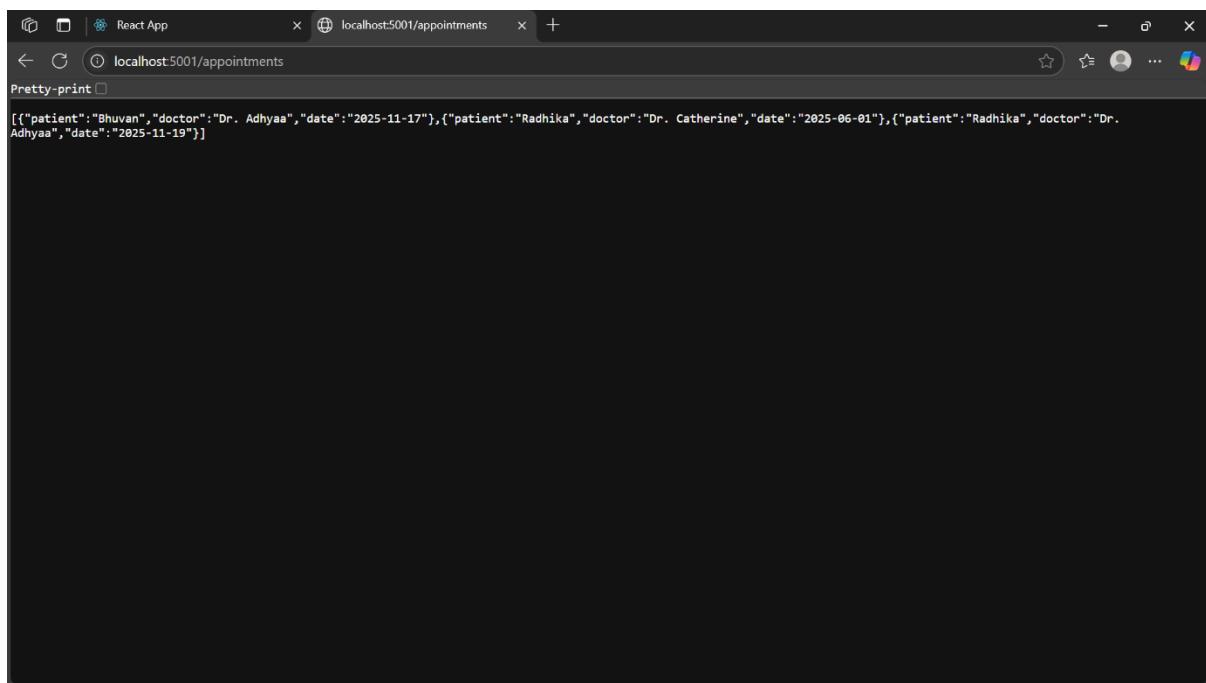
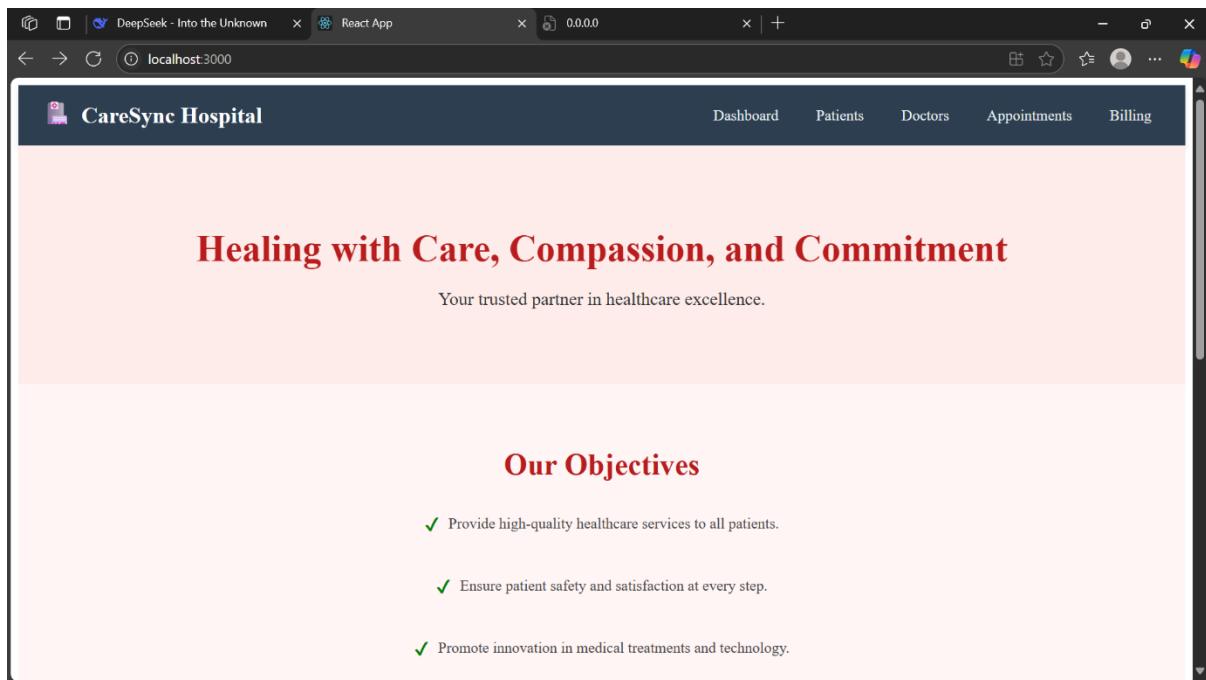
grafana:

image: grafana/grafana

...

All services run inside a shared Docker network (monitoring-net) to ensure proper communication.





Appointments

Select Patient

Select Doctor

dd-mm-yyyy

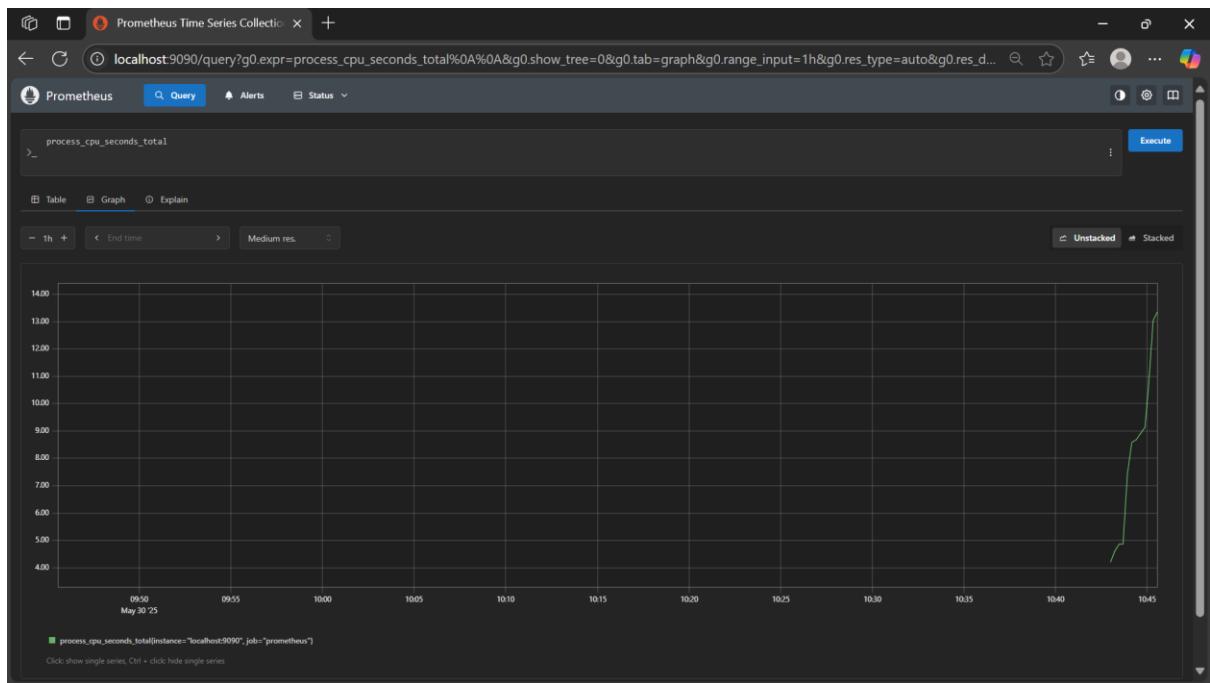
Schedule Appointment

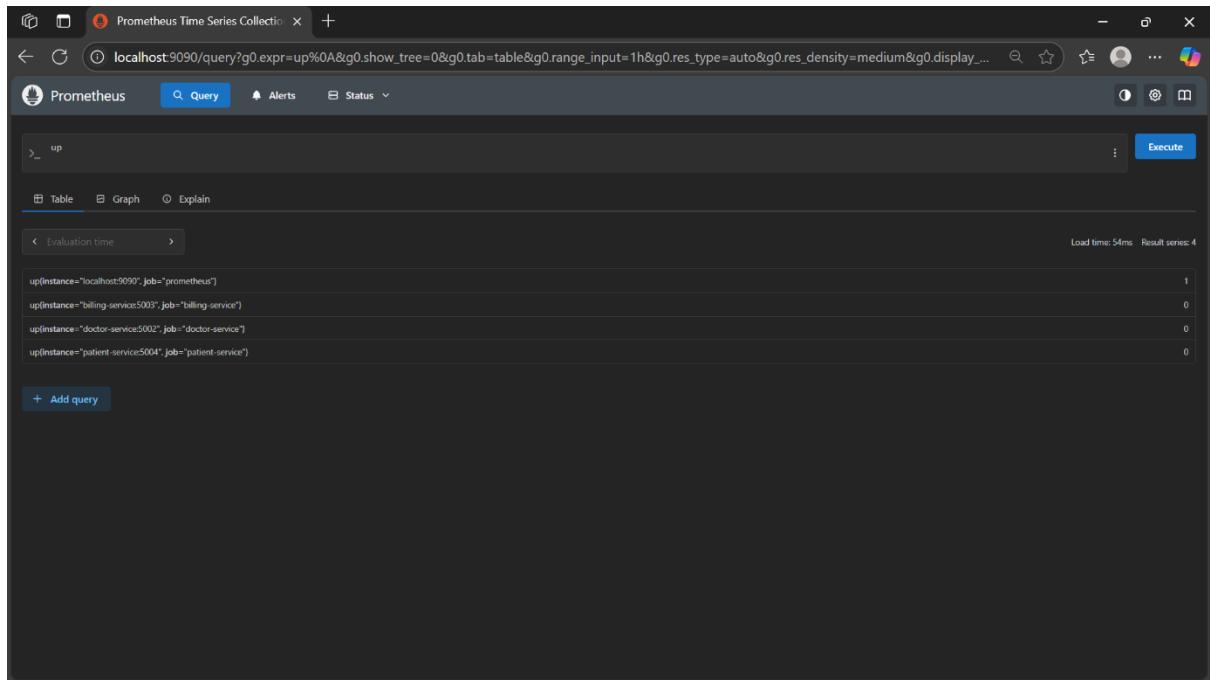
Patient: Bhuvan
Doctor: Dr. Adhyaa
Date: 2025-11-17

Patient:
Doctor:
Date:

Patient:
Doctor:
Date:

PROMETHEUS VISUALIZATION ALONG WITH LOAD BALANCER:

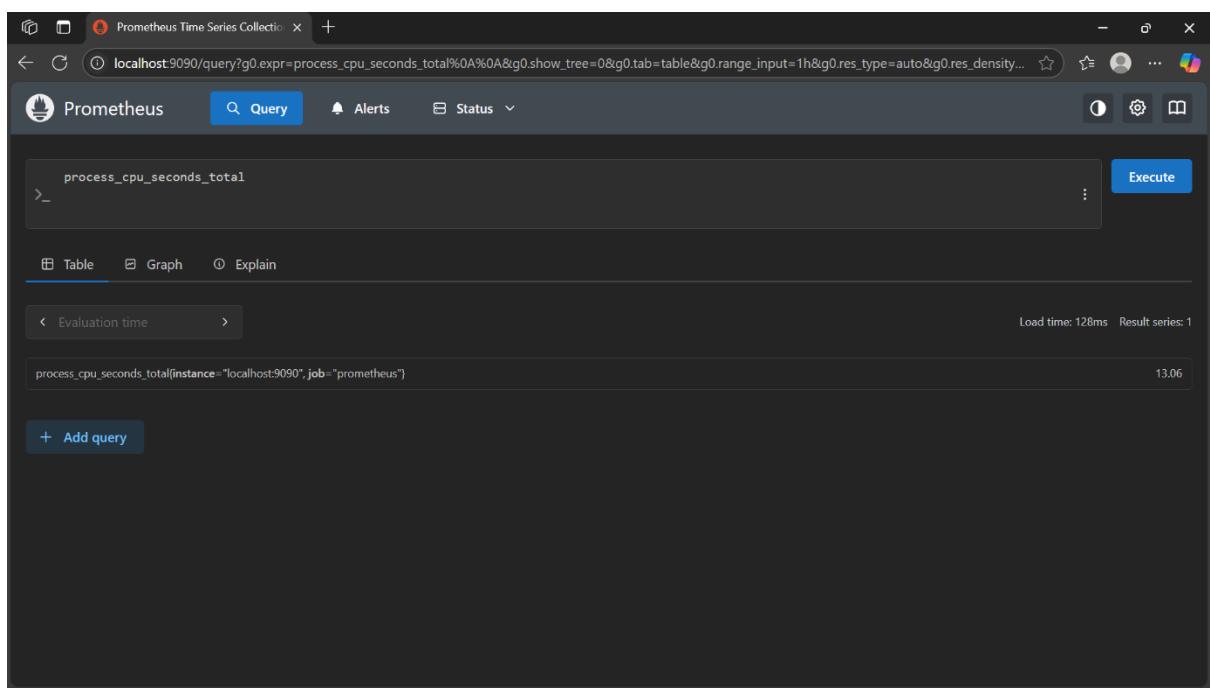




The screenshot shows the Prometheus web interface with a dark theme. The top navigation bar includes tabs for 'Query', 'Alerts', and 'Status'. The main area displays a table of results for the query `up`. The table has three columns: 'Metric' (showing the full query), 'Value' (showing the result), and 'Label' (which is empty). The results are:

Metric	Value	Label
<code>up{instance="localhost:9090", job="prometheus"}</code>	1	
<code>up{instance="billing-service:5003", job="billing-service"}</code>	0	
<code>up{instance="doctor-service:5002", job="doctor-service"}</code>	0	
<code>up{instance="patient-service:5004", job="patient-service"}</code>	0	

Below the table, there is a button labeled '+ Add query'.



The screenshot shows the Prometheus web interface with a dark theme. The top navigation bar includes tabs for 'Query', 'Alerts', and 'Status'. The main area displays a table of results for the query `process_cpu_seconds_total`. The table has three columns: 'Metric' (showing the full query), 'Value' (showing the result), and 'Label' (which is empty). The results are:

Metric	Value	Label
<code>process_cpu_seconds_total{instance="localhost:9090", job="prometheus"}</code>	13.06	

Below the table, there is a button labeled '+ Add query'.

Lab Evaluation 5: Configuration Management Using Ansible

What is Ansible?

Ansible is an **open-source automation tool** used for:

- **Provisioning** servers (setting them up)
- **Configuring** software/services
- **Deploying applications**
- **Managing multiple machines** at once

It follows the **Infrastructure as Code (IaC)** approach and is **agentless**, meaning it does not require any software to be installed on the target machines.

Ansible Architecture Overview

- **Control Node:** The machine where Ansible is installed (usually your local system or admin server).
- **Managed Nodes:** The servers you want to configure using Ansible (e.g., Ubuntu servers).
- **Inventory File:** A file listing the IP addresses or hostnames of all managed nodes.
- **Playbook:** A YAML file that defines the automation tasks to be run on the nodes.

Steps to Configure and Use Ansible

Step 1: Install Ansible (on Control Node)

For Ubuntu/Debian:

```
bash
```

```
sudo apt update  
sudo apt install ansible -y
```

To verify:

```
bash  
ansible --version
```

Step 2: Set Up SSH Access to Managed Nodes

- Ensure the control node can connect to managed nodes via **SSH**.

```
bash  
ssh-copy-id user@<managed-node-ip>
```

This command adds your control node's SSH key to the authorized keys of the managed node so that it can connect without a password.

Step 3: Create an Inventory File

The **inventory file** tells Ansible which servers to manage.

Example (hosts file):

```
ini  
[webservers]  
192.168.1.101  
192.168.1.102
```

```
[dbservers]  
192.168.1.201
```

Save this file as hosts or inventory.txt.

Step 4: Ping Test to Check Connectivity

Use Ansible's built-in ping module to test the connection:

bash

```
ansible all -i hosts -m ping
```

Expected output:

ruby

```
192.168.1.101 | SUCCESS => {
```

```
    "changed": false,  
    "ping": "pong"  
}
```

Step 5: Write Your First Ansible Playbook

A **playbook** is a YAML file that defines what tasks to perform.

Example (install_apache.yaml):

yaml

```
---
```

```
- name: Install Apache on webservers  
  hosts: webservers  
  become: true # Runs tasks as root (sudo)
```

tasks:

```
  - name: Install Apache package
```

```
apt:  
  name: apache2  
  state: present  
  update_cache: yes  
  
- name: Start Apache service  
  service:  
    name: apache2  
    state: started  
    enabled: yes
```

Step 6: Run the Playbook

Use the following command:

bash

```
ansible-playbook -i hosts install_apache.yaml
```

This will:

- Install Apache on all servers under [webservers]
- Ensure the service is started and enabled at boot

Step 7: Use Roles for Structured Playbooks (Optional but Recommended)

Roles allow us to organize our playbooks in a reusable way.

Folder structure:

css

```
roles/  
  apache/  
    tasks/  
      main.yaml  
    handlers/  
      main.yaml  
  templates/  
  files/
```

Then your main playbook includes:

```
yaml  
- hosts: webservers  
  roles:  
    - apache
```

Step 8: Verify and Debug

- Use ansible-playbook with -v or -vvv for verbose output.
- Check logs and execution results directly from terminal.

Objective:

- Automate infrastructure provisioning and system configuration using Ansible.
- Apply Infrastructure as Code (IaC) concepts for cloud environments.

Tasks Performed:

- Installed and configured Ansible on a control node.
- Defined inventory files to manage multiple target hosts.
- Wrote Ansible playbooks to automate tasks like installing packages, starting services, and configuring files.
- Tested idempotency and error handling of playbooks.
- Simulated deployment of a web server stack using Ansible roles.

Tools/Technologies Used:

- Ansible CLI
- YAML (for writing playbooks)
- Ubuntu/Linux servers (as managed nodes)
- SSH (for remote execution)

Outcome:

- Developed automation scripts to streamline infrastructure management.
- Demonstrated knowledge of Ansible architecture, inventory management, and playbooks.
- Enhanced deployment speed, consistency, and reduced human error in system configuration.
- Laid the foundation for implementing full DevOps pipelines in cloud environments.

```
rakshayani@LAPTOP-HDHT03IO:~  
sudo: systemctl: command not found  
rakshayani@LAPTOP-HDHT03IO:~$ sudo systemctl start docker  
rakshayani@LAPTOP-HDHT03IO:~$ sudo systemctl enable docker  
rakshayani@LAPTOP-HDHT03IO:~$ docker --version  
Docker: version 26.1.3, build 26.1.3-0ubuntu1=24.04.1  
rakshayani@LAPTOP-HDHT03IO:~$ sudo systemctl status docker  
● docker.service - Docker Application Container Engine  
   Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; preset: enabled)  
     Active: active (running) since Thu 2025-05-29 14:45:46 UTC; 2min 46s ago  
TriggeredBy: ● docker.socket  
   Docs: https://docs.docker.com  
 Main PID: 1693 (dockerd)  
    Tasks: 13  
   Memory: 54.9M ()  
    CGroup: /system.slice/docker.service  
           └─1693 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock  
  
May 29 14:45:40 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:40.653746190Z" level=info msg=""  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.771876146Z" level=info msg="  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950092345Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950124961Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950129500Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950129500Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950129500Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950169977Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950169977Z" level=warning >  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.970963201Z" level=info msg=""  
May 29 14:45:46 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:46.181751289Z" level=info msg=""  
May 29 14:45:46 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:46.181751289Z" level=info msg=""  
May 29 14:45:46 LAPTOP-HDHT03IO systemd[1]: Started docker.service - Docker Application Container Engine  
lines 1-21/21 (END), skipping...  
● docker.service - Docker Application Container Engine  
   Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; preset: enabled)  
     Active: active (running) since Thu 2025-05-29 14:45:46 UTC; 2min 46s ago  
TriggeredBy: ● docker.socket  
   Docs: https://docs.docker.com  
 Main PID: 1693 (dockerd)  
    Tasks: 13  
   Memory: 54.9M ()  
    CGroup: /system.slice/docker.service  
           └─1693 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/containerd.sock  
  
May 29 14:45:40 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:40.653746190Z" level=info msg="Loading containers: start."  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.771876146Z" level=info msg="Loading containers: done."  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950092345Z" level=warning msg="WARNING: No blkio throttle.read_bps device support"  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950124961Z" level=warning msg="WARNING: No blkio throttle.write_bps device support"  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950129500Z" level=warning msg="WARNING: No blkio throttle.read_iops_device support"  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950129500Z" level=warning msg="WARNING: No blkio throttle.write_iops_device support"  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.950169977Z" level=info msg="Docker daemon" commit="26.1.3-0ubuntu1=24.04.1" containerd-snapshotter=false storage-driver=overlay2 version=>  
May 29 14:45:42 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:42.970963201Z" level=info msg="Daemon has completed initialization"  
May 29 14:45:46 LAPTOP-HDHT03IO dockerd[1693]: time="2025-05-29T14:45:46.181751289Z" level=info msg="API listen on /run/docker.sock"  
May 29 14:45:46 LAPTOP-HDHT03IO systemd[1]: Started docker.service - Docker Application Container Engine.
```